CAS CS 538. Solutions to Problem Set 2

Due electronically via gradescope, Monday February 2, 2026 11:59pm

Useful Definitions

A function is negligible is if goes to zero very quickly. How quickly? Faster than $1/n$, $1/n^2$, $1/n^3$, etc. Formally, we have the following definition.

Definition 1 (Negligible function). A function $f : \mathbb{Z}_{\geqslant 1} \to \mathbb{R}$ is **negligible** if for all integers $c > 0$, we have

$$\lim_{n \to \infty} f(n) \cdot n^c = 0.$$

We will need negligible functions to formalize security (basically, we will want to make sure that the probability of adversarial success is negligible). See Chapter 2.3.1 of Boneh & Shoup textbook.

Definition 2 (Semantic security advantage). For a given cipher $\mathcal{E} = (E, D)$ defined over $(\mathcal{K}, \mathcal{M}\mathcal{C})$ and for a given adversary $\mathcal{A}$, we define two experiments, Experiment 0 and Experiment 1. For $b \in \{0, 1\}$, we define Experiment b as follows:

- The adversary computes $m_0, m_1 \in \mathcal{M}$, of the same length, and sends them to the challenger.

- The challenger computs $k \xleftarrow{\text{R}} \mathcal{K}, c \xleftarrow{\text{R}} E(k, m_b)$, and sends c to the adversary.

- The adversary outputs a bit $\hat{b} \in \{0, 1\}$.

For $b \in \{0, 1\}$, let W_b be the event that $\mathcal{A}$ outputs 1 in Experiment b. We define $\mathcal{A}$'s **semantic security advantage** with respect to $\mathcal{E}$ as

$$\text{SSadv}[\mathcal{A}, \mathcal{E}] := |\Pr[W_0] - \Pr[W_1]|.$$

Definition 3 (Semantic security). A cipher $\mathcal{E}$ is **semantically secure** if for all efficient adversaries $\mathcal{A}$, the value $\text{SSadv}[\mathcal{A}, \mathcal{E}]$ is negligible.

Problems

Problem 1. (15 points)
   Show that each of the following functions are negligible:

   a) $f(n) + g(n)$, where $f(n)$ and $g(n)$ are both negligible

   b) $1000 \cdot f(n)$, where $f(n)$ is negligible

   c) $\frac{1}{2^{\sqrt{n}}}$

Solution. For **a**, let $c > 0$ be any positive int. Since $f(n)$ is negligible, by its definition we have:

$$\lim_{n \to \inf} f(n) \cdot n^c = 0$$

The same can be said for $g(n)$ too:

$$\lim_{n \to \inf} g(n) \cdot n^c = 0$$

Using limit properties (property of sums), we can then get:

$$\lim_{n \to \infty} [f(n) + g(n)] \cdot n^c = \lim_{n \to \infty} [f(n) \cdot n^c + g(n) \cdot n^c]$$
$$= \lim_{n \to \infty} f(n) \cdot n^c + \lim_{n \to \infty} g(n) \cdot n^c$$
$$= 0 + 0 = 0$$

Since this holds for any int $c > 0$, this proves that **a** is negligible.

For **b**, let $c > 0$ be any positive int. Like before, $f(n)$ is negligible, so:

$$\lim_{n \to \inf} f(n) \cdot n^c = 0$$

Using limit properties (that of multiples), we then get:

$$\lim_{n \to \infty} [1000 \cdot f(n)] \cdot n^c = \lim_{n \to \infty} 1000 \cdot [f(n) \cdot n^c]$$
$$= 1000 \cdot \lim_{n \to \infty} f(n) \cdot n^c$$
$$= 1000 \cdot 0 = 0$$

As it holds for all $c > 0$, this proves **b** is negligible.

For **c**, let $c > 0$ be any positive int. We now need to show that, as n grows, $\frac{1}{2^{\sqrt{n}}}$ will decay into something negligible (so to 0 essentially). Firstly, consider the expression:

$$\frac{1}{2^{\sqrt{n}}} \cdot n^c = \frac{n^c}{2^{\sqrt{n}}}$$

I'll show now how $\lim_{n \to \infty} \frac{n^c}{2^{\sqrt{n}}} = 0$ using limit properties.
Firstly, take the natural log:

$$\ln(\frac{n^c}{2^{\sqrt{n}}}) = \ln(n^c) - \ln(2^{\sqrt{n}}) = c \ln n - \sqrt{n} \ln 2$$

Taking this equation, we can now consider it as a limit as $n \to \infty$:

$$\lim_{n \to \infty} [c \ln n - \sqrt{n} \ln 2] = ???$$

Now looking at this expression, consider *which* term would grow faster consider the circumstances. $c \ln n$, or $\sqrt{n} \ln 2$? Well, the first term is essentially being multipled by a constant, c. The latter term, however, **is the multiple**. Therefore we can assume that the 2nd term will dominate this growth.

$$\lim_{n \to \infty} [c \ln n - \sqrt{n} \ln 2] = -\infty$$

Now we can take this result and plug it back into our original negligible function that we're considering:

$$\lim_{n \to \infty} \frac{n^c}{2^{\sqrt{n}}} = \lim_{n \to \infty} e^{\ln(\frac{n^c}{2^{\sqrt{n}}})}$$
$$= e^{-\infty}$$
$$= 0$$

And since this applies for all $c > 0$, this proves **c**.

Problem 2. (30 points)

You're a rising spy working your way up the ranks of national intelligence. One day you are sent a top-secret mission via a mysterious encrypted message. It's a single clue: a three-character airport code for one of the following airports in five possible countries:

USA	Germany	China	Brazil	India
BOS	MUC	PEK	GIG	DEL
JFK	FRA	SZX	BSB	BOM
LAX	DUS	PVG	GRU	MAA
DCA	HAM	CAN	CGH	BLR

Unbeknownst to you, a copy of the ciphertext was intercepted by a rival attempting to learn which country you are going to.

Consider the following security game between a challenger and an adversary. Let $\Sigma = (\mathsf{Enc}, \mathsf{Dec})$ be a computational cipher with message space $\mathcal{M}$ equal to the set of 20 airport codes above. The challenger chooses a random plaintext $m \xleftarrow{\text{R}} \mathcal{M}$ and key $k \xleftarrow{\text{R}} \mathcal{K}$, computes $c \xleftarrow{\text{R}} \mathsf{Enc}(k, m)$ and sends c to the adversary. The adversary $\mathcal{A}$ outputs a string s from the following set:

$$\{\texttt{"USA"}, \texttt{"Germany"}, \texttt{"China"}, \texttt{"Brazil"}, \texttt{"India"}\}.$$

Let $\mathsf{Country}(m)$ be a function that outputs the country string in which a given airport code is located. Let W be the event that $s = \mathsf{Country}(m)$, let $p = \Pr[W]$, and define $\mathsf{CountryGuessAdv}[\mathcal{A}, \Sigma] = |p - \frac{1}{5}|$.

Prove that if Σ is a semantically secure cipher (per Section 2.2.2 of the textbook), then for any efficient adversary $\mathcal{A}$, $\mathsf{CountryGuessAdv}[\mathcal{A}, \Sigma]$ is negligible. Notice what this implies: your rival's attempt to learn which country you're travelling to is at most negligibly better than random guessing.

To prove this, use a reduction: supposing there exists an efficient adversary with non-negligible advantage in the country-guessing game, construct an efficient adversary with non-negligible advantage in the semantic security game.

Solution. I'll prove the **contrapositive** here, so: If there exists some efficent adversary $\mathcal{A}$ with a non-negligible advantage, then we can make an efficent adversary $\mathcal{B}$ also with a non-negligible advantage in the semantic security game, meaning that Σ isn't secure.

Given $\mathcal{A}$ for the country-guessing game with CountryGuessAdv$[\mathcal{A}, \Sigma] = \epsilon$, we can make $\mathcal{B}$ for the semantic security game as follows:

- $\mathcal{B}$ will choose 2 distinct countries C_0, C_1 from USA, Germany, China, Brazil, India.

- For each country C_i, $\mathcal{B}$ choose some arbitary airport code m_i from that country.

- $\mathcal{B}$ will then submit m_0, m_1 to the semantic security challenger (*important to note how these codes will have the same length, all codes are 3-characters long.*)

- The challenger chooses $b \leftarrow 0, 1$, $k \leftarrow \mathcal{K}$, computs $c \leftarrow \mathsf{Enc}(k, m_b)$, and then sends c to $\mathcal{B}$.

- $\mathcal{B}$ then sends result c to $\mathcal{A}$ and recieves $\mathcal{A}$'s output
  $s \in$ "USA", "Germany", "China", "Brazil", "India".

- $\mathcal{B}$ outputs $\hat{b}$ in accordance to the following rule:

  - If $s = \mathsf{Country}(m_0)$, output $\hat{b} = 0$.
  - If $s = \mathsf{Country}(m_1)$, output $\hat{b} = 1$
  - Elser, output rnd $\hat{b} \leftarrow 0, 1$

Let W_b be the occurance that $\mathcal{B}$ outputs 1 in Experiement b of the semantic security game.

Case 1: $b = 0$ (challenger encrypts m_0)

- This case basically says that $\mathcal{A}$ recieves m_0 as encryption.

- Let $p_0 = \Pr[\mathcal{A}$ outputs $\mathsf{Country}(m_0) \mid$ encryption of $m_0]$.

- Let $p_1 = \Pr[\mathcal{A}$ outputs $\mathsf{Country}(m_1) \mid$ encryption of $m_0]$.

- Then, $\Pr[W_0] = \Pr[\mathcal{B}$ outputs $1 \mid b = 0] = p_1 + \frac{1}{2}(1 - p_0 - p_1)$

Case 2: $b = 1$ (challenger encrpyts m_1)

- In this case, $\mathcal{A}$ insteads gets m_1.

- Let $q_0 = \Pr[\mathcal{A}$ outputs $\mathsf{Country}(m_0) \mid$ encryption of $m_1]$.

- Let $q_1 = \Pr[\mathcal{A}$ outputs $\mathsf{Country}(m_1) \mid$ encryption of $m_1]$.

- Then, $\Pr[W_1] = \Pr[\mathcal{B}$ outputs $1 \mid b = 1] = q_1 + \frac{1}{2}(1 - q_0 - q_1)$.

By the definition of CountryGuessAdv, we know that, overall, $\mathcal{A}$ will guess the correct country with a probability of $\frac{1}{5} + \delta$ for some δ (where $|\delta| = \epsilon$, and we can then also assume that *,in general*, that $\delta > 0$ because otherwise we could just flip $\mathcal{A}$'s output).

When looking at the specific countries C_0, C_1, $\mathcal{A}$ NEEDS to have AT LEAST the avg. adv. Then *in general*, we then assume that $\mathcal{A}$ is AT LEAST as good at distinguishing between C_0, C_1 as between any other pair of countries. We then get:

$$\frac{p_0 + q_1}{2} \geqslant \frac{1}{5} + \frac{\epsilon}{2}$$

This is because, when encryption m_0, $\mathcal{A}$ should output $\mathsf{Country}(m_0)$, with a probability of AT LEAST $\frac{1}{5} + \frac{\epsilon}{2}$, and similarily of that for m_1.

Then, the semantic security advantage for $\mathcal{B}$ becomes:

$$\begin{aligned}
\mathrm{SSadv}[\mathcal{B}, \Sigma] &= |\Pr[W_0] - \Pr[W_1]| \\
&= \left| \left(p_1 + \frac{1}{2}(1 - p_0 - p_1) \right) - \left(q_1 + \frac{1}{2}(1 - q_0 - q_1) \right) \right| \\
&= \left| \frac{1}{2}(p_1 - p_0 - p_1 + q_0 + q_1 - q_1) \right| \\
&= \frac{1}{2}|q_0 - p_0|
\end{aligned}$$

Now, since $\mathcal{A}$ has that advantage ϵ overall, and we've chosen the counties where $\mathcal{A}$ performs best, we now have $|q_0 - p_0| \geqslant \frac{\epsilon}{5}$. *Note that the factor comes from the num. of possible counties. This means, worst case, $\mathcal{A}$'s adv. is spread evenly across all country pairs.*

Therefore:

$$\mathrm{SSadv}[\mathcal{B}, \Sigma] \geqslant \frac{\epsilon}{10}$$

Since ϵ is non-negligible by our assumption, $\frac{\epsilon}{10}$ will then also be non-negligible (this is proven from our answer from **1** as well, we can apply limit properties to preserve non-negligibility).

By proving this, we ultimately show that $\mathcal{B}$ is an efficent adversary. This proves the contrapositive of the OG statements, therefore proving the original statement.

Problem 3. (30 points) Let Σ be a computational cipher with $|\mathcal{K}| < |\mathcal{M}|$. Construct an adversary $\mathcal{A}$ against Σ such that the running time of $\mathcal{A}$ is very reasonable (in fact, comparable to the running time of Enc and Dec) and $\mathrm{SSadv}[\mathcal{A}, \Sigma] > 0$. Note that it is okay if $\mathrm{SSadv}[\mathcal{A}, \Sigma]$ is very small, as long as it is positive.

You must demonstrate what the adversary does and prove that its SSadv is positive. The adversary should **not** depend on any knowledge about Σ that cannot be efficiently obtained (for example, the adversary doesn't know exact probabilities of different ciphertexts).

Hint: For the adversary design, use the idea from Discussion 2 Problem 2. You can't perform exhaustive search, of course, because you don't have the time; make a random guess instead. Then analyze the probability of outputting 1 in each of the two experiments. This analysis will be different from the one in discussion — go

back to your first principles of probability. Conclude that it has non-zero advantage.

This problem justifies why any reasonable definition of semantic security must allow for at least a negligible advantage for $\mathcal{A}$ even when the running time of $\mathcal{A}$ is limited.

Solution. Back in last week's discussion, the solution to the referred question involved us essentially doing this exhaustive key search, eventually getting us that SSadv, but being really ineffecient. This attack however, could just be accomplished by basically *randomly choosing msgs and keys and interpreting their ciphertexts and outputs.*

Given the cipher with the above properties, we can make an adversary $\mathcal{A}$ as follow:

- $\mathcal{A}$ chooses 2 distinct msgs $m_0, m_1 \leftarrow \mathcal{M}$ uniformly at rnd (making sure $m_0 \neq m_1$).

- $\mathcal{A}$ sends m_0, m_1 to the challenger.

- The challenger outputs ciphertext $c \leftarrow \mathsf{Enc}(k, m_b)$ for $b \leftarrow 0, 1$.

- $\mathcal{A}$ chooses a rnd key $k' \leftarrow \mathcal{K}$ uniformly.

- $\mathcal{A}$ computes $m' \leftarrow \mathsf{Dec}(k', c)$.

- $\mathcal{A}$ outputs 1 if $m' = m_0$, and 0 otherwise.

The running time of this adversary is dominated by the one Dec operation, and is therefore comparable to the running time of Enc and Dec. This satisfies the time constraint for this problem.

Now, let W_b be the occurance that $\mathcal{A}$ outputs 1 in Experiment b. I can now analyze $\Pr[W_0]$ and $\Pr[W_1]$.
Experiment 1: $b = 0$ (c is an encryption of m_0 with some key k^*)

- The real key k^* satisfies $\mathsf{Dec}(k^*, c) = m_0$

- When $\mathcal{A}$ picks $k' = k^*$ (which happens with a chance of $1/|\mathcal{K}|$), $\mathsf{Dec}(k', c) = m_0$

- When $\mathcal{A}$ picks $k' \neq k^*$, $\mathsf{Dec}(k', c)$ could be m_0 or something else

- Let $p = \Pr[\mathsf{Dec}(k', c) = m_0 \mid k' \neq k^*]$

- Then,

$$\Pr[W_0] = \frac{1}{|\mathcal{K}|} + (1 - \frac{1}{|\mathcal{K}|})p$$

Experiment 2: $b = 1$ (c is an encryption of m_1 with some key k^*.)

- Now, the real key k^* satisfies $\mathsf{Dec}(k^*, c) = m_1 \neq m_0$

- When $\mathcal{A}$ picks $k' = k^*$ (the probability is $1/|\mathcal{K}|$), $\mathsf{Dec}(k', c) = m_1 \neq m_0$

- When $\mathcal{A}$ picks $k' \neq k^*$, $\mathsf{Dec}(k', c)$ can by any other wrong answer.

- Let $q = \Pr[\mathrm{Dec}(k', c) = m_0 \mid k' \neq k^*]$

- Therefore,

$$\Pr[W_1] = 0 \cdot \frac{1}{|\mathcal{K}|} + (1 - \frac{1}{|\mathcal{K}|})q = (1 - \frac{1}{|\mathcal{K}|})q$$

Now, since $|\mathcal{K}| < |\mathcal{M}|$, each ciphertext c can be decrypted to AT MOST $|\mathcal{K}|$ different messages (one / key). Over all possible $k' \in \mathcal{K}$, the values $\mathrm{Dec}(k', c)$ take AT MOST $|\mathcal{K}|$ distinct values. Since m_0 was chosen uniformly from $\mathcal{M}$ and $|\mathcal{K}| < |\mathcal{M}|$:

$$q \leqslant \frac{|\mathcal{K}|}{|\mathcal{M}|} < 1$$

Moreover, experiment 0 shows show that we know AT LEAST 1 key (k^*) that does decrypt to m_0, so:

$$p \geqslant \frac{1}{|\mathcal{K}|}(\text{or, the fraction of keys that decrypt } c \text{ to } m_0)$$

This can also be thought of as:

$$p = \frac{\text{number of keys } k' \text{ with } \mathrm{Dec}(k', c) = m_0}{|\mathcal{K}|} \geqslant \frac{1}{|\mathcal{K}|}$$

And therefore:

$$\Pr[W_0] - \Pr[W_1]$$
$$= [\frac{1}{|\mathcal{K}|} + (1 - \frac{1}{|\mathcal{K}|})p] - (1 - \frac{1}{|\mathcal{K}|})q$$
$$= \frac{1}{|\mathcal{K}|} + (1 - \frac{1}{|\mathcal{K}|})(p - q)$$

Since $p \geqslant 1/|\mathcal{K}|$ and $q \leqslant |\mathcal{K}|/|\mathcal{M}| < 1$, and $1/|\mathcal{K}| > 0$, we have:

$$\Pr[W_0] - \Pr[W_1] > 0$$

Thusly, $\mathrm{SSadv}[\mathcal{A}, \Sigma] = |\Pr[W_0] - \Pr[W_1]| > 0$, giving us an efficent attack that breaks the original cipher algorithm.

Problem 4. (25 points)[Boneh-Shoup Exercise 2.10 from Section 2.6] Let $\Sigma = (\mathrm{Enc}, \mathrm{Dec})$ be a semantically secure cipher defined over $(\mathcal{K}, \mathcal{M}, \mathcal{C})$, where $\mathcal{M} = \mathcal{C} = \{0, 1\}^L$. Which of the following encryption algorithms yields a semantically secure scheme? Either give an attack or provide a security proof via an explicit reduction.

(a) Let Σ_1 be a scheme s.t. $\mathrm{Enc}_1(k, m) := 0 || \mathrm{Enc}(k, m)$

Solution. Your solution goes here

(b) Let Σ_2 be a scheme s.t. $\mathrm{Enc}_2(k, m) := \mathrm{Enc}(k, m) || \mathrm{parity}(m)$, where parity of a binary string refers to the number of 1 bits (equivalently, the exclusive-or of all the bits) in the string.

Solution. Your solution goes here