

# Chapter 3

## Stream ciphers

In the previous chapter, we introduced the notions of perfectly secure encryption and semantically secure encryption. The problem with perfect security is that to achieve it, one must use very long keys. Semantic security was introduced as a weaker notion of security that would perhaps allow us to build secure ciphers that use reasonably short keys; however, we have not yet produced any such ciphers. This chapter studies one type of cipher that does this: the stream cipher.

### 3.1 Pseudo-random generators

Recall the one-time pad. Here, keys, messages, and ciphertexts are all  $L$ -bit strings. However, we would like to use a key that is much shorter. So the idea is to instead use a short,  $\ell$ -bit “seed”  $s$  as the encryption key, where  $\ell$  is much smaller than  $L$ , and to “stretch” this seed into a longer,  $L$ -bit string that is used to mask the message (and unmask the ciphertext). The string  $s$  is stretched using some efficient, deterministic algorithm  $G$  that maps  $\ell$ -bit strings to  $L$ -bit strings. Thus, the key space for this modified one-time pad is  $\{0, 1\}^\ell$ , while the message and ciphertext spaces are  $\{0, 1\}^L$ . For  $s \in \{0, 1\}^\ell$  and  $m, c \in \{0, 1\}^L$ , encryption and decryption are defined as follows:

$$E(s, m) := G(s) \oplus m \quad \text{and} \quad D(s, c) := G(s) \oplus c.$$

This modified one-time pad is called a **stream cipher**, and the function  $G$  is called a **pseudo-random generator**.

If  $\ell < L$ , then by Shannon’s Theorem, this stream cipher cannot achieve perfect security; however, if  $G$  satisfies an appropriate security property, then this cipher is semantically secure. Suppose  $s$  is a random  $\ell$ -bit string and  $r$  is a random  $L$ -bit string. Intuitively, if an adversary cannot effectively tell the difference between  $G(s)$  and  $r$ , then he should not be able to tell the difference between this stream cipher and a one-time pad; moreover, since the latter cipher is semantically secure, so should be the former. To make this reasoning rigorous, we need to formalize the notion that an adversary cannot “effectively tell the difference between  $G(s)$  and  $r$ .”

An algorithm that is used to distinguish a pseudo-random string  $G(s)$  from a truly random string  $r$  is called a **statistical test**. It takes a string as input, and outputs 0 or 1. Such a test is called **effective** if the probability that it outputs 1 on a pseudo-random input is significantly different than the probability that it outputs 1 on a truly random input. Even a relatively small difference in probabilities, say 1%, is considered significant; indeed, even with a 1% difference, if we can obtain a few hundred independent samples, which are either all pseudo-random or all truly

random, then we will be able to infer with high confidence whether we are looking at pseudo-random strings or at truly random strings. However, a non-zero but negligible difference in probabilities, say  $2^{-100}$ , is not helpful.

How might one go about designing an effective statistical test? One basic approach is the following: given an  $L$ -bit string, calculate some statistic, and then see if this statistic differs greatly from what one would expect if the string were truly random.

For example, a very simple statistic that is easy to compute is the number  $k$  of 1's appearing in the string. For a truly random string, we would expect  $k \approx L/2$ . If the PRG  $G$  had some bias towards either 0-bits or 1-bits, we could effectively detect this with a statistical test that, say, outputs 1 if  $|k - 0.5L| < 0.01L$ , and otherwise outputs 0. This statistical test would be quite effective if the PRG  $G$  did indeed have some significant bias towards either 0 or 1.

The test in the previous example can be strengthened by considering not just individual bits, but pairs of bits. One could break the  $L$ -bit string up into  $\approx L/2$  bit pairs, and count the number  $k_{00}$  of pairs 00, the number  $k_{01}$  of pairs 01, the number  $k_{10}$  of pairs 10, and the number  $k_{11}$  of pairs 11. For a truly random string, one would expect each of these numbers to be  $\approx L/2 \cdot 1/4 = L/8$ . Thus, a natural statistical test would be one that tests if the distance from  $L/8$  of each of these numbers is less than some specified bound. Alternatively, one could sum up the squares of these distances, and test whether this sum is less than some specified bound — this is the classical  $\chi^2$  test from statistics. Obviously, this idea generalizes from pairs of bits to tuples of any length.

There are many other simple statistics one might check. However, simple tests such as these do not tend to exploit deeper mathematical properties of the algorithm  $G$  that a malicious adversary may be able to exploit in designing a statistical test specifically geared towards  $G$ . For example, there are PRG's for which the simple tests in the previous two paragraphs are completely ineffective, but yet are completely predictable, given sufficiently many output bits; that is, given a prefix of  $G(s)$  of sufficient length, the adversary can compute all the remaining bits of  $G(s)$ , or perhaps even compute the seed  $s$  itself.

Our definition of security for a PRG formalizes the notion that there should be no effective (and efficiently computable) statistical test.

### 3.1.1 Definition of a pseudo-random generator

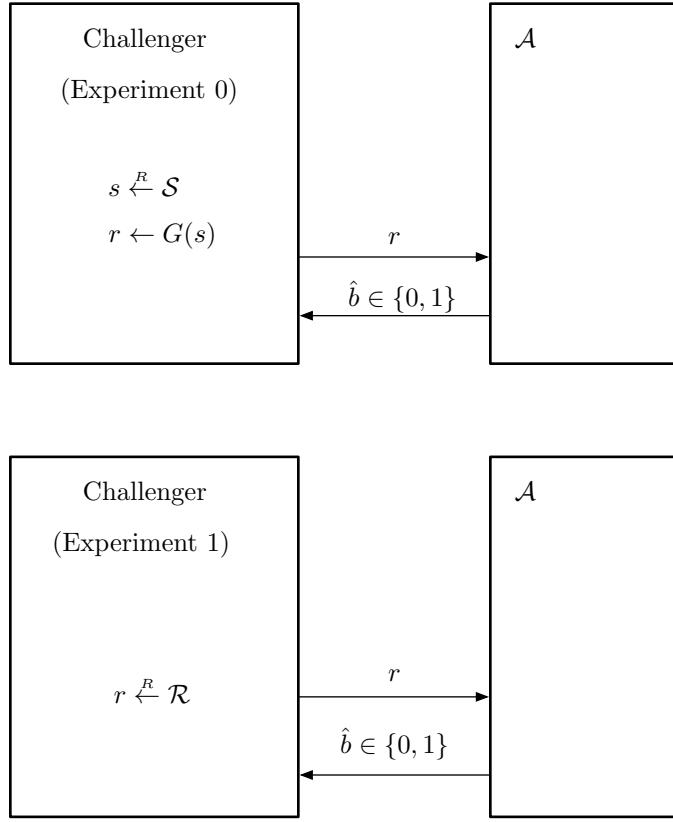
A **pseudo-random generator**, or **PRG** for short, is an efficient, deterministic algorithm  $G$  that, given as input a **seed**  $s$ , computes an output  $r$ . The seed  $s$  comes from a finite **seed space**  $\mathcal{S}$  and the output  $r$  belongs to a finite **output space**  $\mathcal{R}$ . Typically,  $\mathcal{S}$  and  $\mathcal{R}$  are sets of bit strings of some prescribed length (for example, in the discussion above, we had  $\mathcal{S} = \{0, 1\}^\ell$  and  $\mathcal{R} = \{0, 1\}^L$ ). We say that  $G$  is a PRG defined over  $(\mathcal{S}, \mathcal{R})$ .

Our definition of security for a PRG captures the intuitive notion that if  $s$  is chosen at random from  $\mathcal{S}$  and  $r$  is chosen at random from  $\mathcal{R}$ , then no efficient adversary can effectively tell the difference between  $G(s)$  and  $r$ : the two are **computationally indistinguishable**. The definition is formulated as an attack game.

**Attack Game 3.1 (PRG).** For a given PRG  $G$ , defined over  $(\mathcal{S}, \mathcal{R})$ , and for a given adversary  $\mathcal{A}$ , we define two experiments, Experiment 0 and Experiment 1. For  $b = 0, 1$ , we define:

**Experiment  $b$ :**

- The challenger computes  $r \in \mathcal{R}$  as follows:



**Figure 3.1:** Experiments 0 and 1 of Attack Game 3.1

- if  $b = 0$ :  $s \leftarrow^R \mathcal{S}$ ,  $r \leftarrow G(s)$ ;
- if  $b = 1$ :  $r \leftarrow^R \mathcal{R}$ .

and sends  $r$  to the adversary.

- Given  $r$ , the adversary computes and outputs a bit  $\hat{b} \in \{0, 1\}$ .

For  $b = 0, 1$ , let  $W_b$  be the event that  $\mathcal{A}$  outputs 1 in Experiment  $b$ . We define  $\mathcal{A}$ 's **advantage** with respect to  $G$  as

$$\text{PRGadv}[\mathcal{A}, G] := \left| \Pr[W_0] - \Pr[W_1] \right|. \quad \square$$

The attack game is illustrated in Fig. 3.1.

**Definition 3.1 (secure PRG).** A PRG  $G$  is **secure** if the value  $\text{PRGadv}[\mathcal{A}, G]$  is negligible for all efficient adversaries  $\mathcal{A}$ .

As discussed in Section 2.2.5, Attack Game 3.1 can be recast as a “bit guessing” game, where instead of having two separate experiments, the challenger chooses  $b \in \{0, 1\}$  at random, and then runs Experiment  $b$  against the adversary  $\mathcal{A}$ . In this game, we measure  $\mathcal{A}$ 's *bit-guessing advantage*

$\text{PRGadv}^*[\mathcal{A}, G]$  as  $|\Pr[\hat{b} = b] - 1/2|$ . The general result of Section 2.2.5 (namely, (2.11)) applies here as well:

$$\text{PRGadv}[\mathcal{A}, G] = 2 \cdot \text{PRGadv}^*[\mathcal{A}, G]. \quad (3.1)$$

We also note that a PRG can only be secure if the cardinality of the seed space is super-poly (see Exercise 3.5).

### 3.1.2 Mathematical details

Just as in Section 2.3, we give here more of the mathematical details pertaining to PRGs. Just like Section 2.3, this section may be safely skipped on first reading with very little loss in understanding.

First, we state the precise definition of a PRG, using the terminology introduced in Definition 2.9.

**Definition 3.2 (pseudo-random generator).** A *pseudo-random generator* consists of an algorithm  $G$ , along with two families of spaces with system parameterization  $P$ :

$$\mathbf{S} = \{\mathcal{S}_{\lambda, \Lambda}\}_{\lambda, \Lambda} \quad \text{and} \quad \mathbf{R} = \{\mathcal{R}_{\lambda, \Lambda}\}_{\lambda, \Lambda},$$

such that

1.  $\mathbf{S}$  and  $\mathbf{R}$  are efficiently recognizable and sampleable.
2. Algorithm  $G$  is an efficient deterministic algorithm that on input  $\lambda, \Lambda, s$ , where  $\lambda \in \mathbb{Z}_{\geq 1}$ ,  $\Lambda \in \text{Supp}(P(\lambda))$ , and  $s \in \mathcal{S}_{\lambda, \Lambda}$ , outputs an element of  $\mathcal{R}_{\lambda, \Lambda}$ .

Next, Definition 3.1 needs to be properly interpreted. First, in Attack Game 3.1, it is to be understood that for each value of the security parameter  $\lambda$ , we get a different probability space, determined by the random choices of the challenger and the random choices of the adversary. Second, the challenger generates a system parameter  $\Lambda$ , and sends this to the adversary at the very start of the game. Third, the advantage  $\text{PRGadv}[\mathcal{A}, G]$  is a function of the security parameter  $\lambda$ , and security means that this function is a negligible function.

## 3.2 Stream ciphers: encryption with a PRG

Let  $G$  be a PRG defined over  $(\{0, 1\}^\ell, \{0, 1\}^L)$ ; that is,  $G$  stretches an  $\ell$ -bit seed to an  $L$ -bit output. The **stream cipher**  $\mathcal{E} = (E, D)$  **constructed from**  $G$  is defined over  $(\{0, 1\}^\ell, \{0, 1\}^{\leq L}, \{0, 1\}^{\leq L})$ ; for  $s \in \{0, 1\}^\ell$  and  $m, c \in \{0, 1\}^{\leq L}$ , encryption and decryption are defined as follows: if  $|m| = v$ , then

$$E(s, m) := G(s)[0 \dots v-1] \oplus m,$$

and if  $|c| = v$ , then

$$D(s, c) := G(s)[0 \dots v-1] \oplus c.$$

As the reader may easily verify, this satisfies our definition of a cipher (in particular, the correctness property is satisfied).

Note that for the purposes of analyzing the semantic security of  $\mathcal{E}$ , the length associated with a message  $m$  in Attack Game 2.1 is the natural length  $|m|$  of  $m$  in bits. Also, note that if  $v$  is much smaller than  $L$ , then for many practical PRGs, it is possible to compute the first  $v$  bits of  $G(s)$  much faster than actually computing all the bits of  $G(s)$  and then truncating.

The main result of this section is the following: