

CAS CS 538. Solutions to Problem Set 5

Due electronically via gradescope, Tuesday February 24, 2026 11:59pm
Om Khadka, U51801771

Problem 1. (50 points)

In Discussion 5 we introduced the Signal symmetric ratchet and we showed that even if a subsequent key is leaked to the adversary, it is hard to recover the previous key. In this problem, we will prove something stronger: even if a subsequent key is leaked (and the previous key has been deleted), past encryptions are still semantically secure.

Specifically, let $\mathcal{E} = (\text{Enc}, \text{Dec})$ be a semantically secure cipher over $(\mathcal{K}, \mathcal{M}, \mathcal{C})$. Let $G : \mathcal{S} \rightarrow \mathcal{S} \times \mathcal{K}$ be a secure PRG such that $G : \{0, 1\}^\ell \rightarrow \{0, 1\}^{2\ell}$. Define the symmetric ratchet scheme as $\text{Enc}^1(s, m) = \text{Enc}(G(s)_2, m)$ and $\text{Dec}^1(s, c) = \text{Dec}(G(s)_2, c)$. Here $G(s)_2$ denotes the second part of the output of $G(s)$, i.e. the key part. We refer you to Figure 1 of Discussion 5 for a diagram of the symmetric ratchet.

At time step n , participants will advance the state of the ratchet by computing $(s_n, k_n) \leftarrow G(s_{n-1})$. Then, the past is erased by deleting the old seed s_{n-1} from memory. The new ciphertext is computed as $c_n = \text{Enc}(k_n, m_n)$.

Prove that \mathcal{E}^1 is *semantically secure* for the previous encryption c_{n-1} even if a subsequent seed and key (s_n, k_n) is leaked to the adversary.

Getting started. To get started, you should consider an adversary \mathcal{A} playing a new game called the SSKeyLeak game against \mathcal{E}^1 . The SSKeyLeak is the same as the Semantic Security Game (Boneh + Shoup Attack Game 2.1) except that in addition to a ciphertext c_{n-1} , the adversary *also* receives a leaked seed and key (s_n, k_n) where s_n, k_n are pseudorandom values produced by G^n with an initial random seed.

In particular, say \mathcal{A} outputs two messages m_0, m_1 . Let Game 0 be when $c_{n-1} = \text{Enc}^1(k_{n-1}, m_0)$ and Game 1 be when $c_{n-1} = \text{Enc}(k_{n-1}, m_1)$. \mathcal{A} receives the tuple (c_{n-1}, s_n, k_n) and outputs a bit \hat{b} . Let W_b be the event that $\hat{b} = 1$ in Game b of the SSKeyLeak game. Then we define the key-leakage advantage as $\text{SSKeyLeakadv}[\mathcal{A}, \mathcal{E}^1] = |\Pr[W_0] - \Pr[W_1]|$.

Building a hybrid argument. Use a hybrid argument to prove that \mathcal{E}^1 is semantically secure against leaked keys by showing that $\text{SSKeyLeakadv}[\mathcal{A}, \mathcal{E}^1]$ is negligible.

Hint: you will have to introduce two hybrid games in between Game 0 and Game 1 of SSKeyLeak game. Then you will give three reductions with wrapper adversaries $\mathcal{B}_1, \mathcal{B}_2$, and \mathcal{B}_3 . In no particular order, two of these reductions yield a PRG distinguisher for G^n and one of them yields an SS adversary for \mathcal{E} .

Solution. The whole idea behind how I'll prove this by essentially splitting up the whole process into a bunch of hybrid arguments, and showing how the sum of all of the reductions end up being negligible.

Proof

Let $G^n : \{0, 1\}^\ell \rightarrow \{0, 1\}^{(n+1)\ell}$ be the n -wise sequential composition of G , outputting $(k_1, s_1, \dots, k_n, s_n)$. Also assume that G is a secure PRG, thereby also meaning that G^n is also secure.

The Hybrid Games

Game 0: The initial real SSKeyLeak game w/ bit $b = 0$.

1. Challenger samples $s_0 \leftarrow \{0, 1\}^\ell$.
2. Calculates $(k_1, s_1, \dots, k_n, s_n) \leftarrow G^n(s_0)$.
3. \mathcal{A} then outputs m_0, m_1 .
4. Challenger computes $c_{n-1} = \text{Enc}(k_{n-1}, m_0)$.
5. Challenger then gives (c_{n-1}, k_n, s_n) to \mathcal{A} .
6. \mathcal{A} returns \hat{b} .

Let $\Pr[W_0]$ denote the probability of \mathcal{A} outputting 1.

Hybrid 1: Same as Game 0, but now the keys and seeds are replaced w/ rnd uniform values.

1. Challenger samples $(k_1, \dots, k_n, s_n) \leftarrow \{0, 1\}^{(n+1)\ell}$ uniformly at random.
2. \mathcal{A} outputs m_0, m_1 .
3. Challenger then Computes $c_{n-1} = \text{Enc}(k_{n-1}, m_0)$.
4. Challenger will then give (c_{n-1}, k_n, s_n) to \mathcal{A} .
5. \mathcal{A} finally returns \hat{b} .

$\Pr[H_1]$ will be the probability of \mathcal{A} outputting 1.

Hybrid 2: Hybrid 1, but now the msg is m_1 .

1. Challenger samples $(k_1, \dots, k_n, s_n) \leftarrow \{0, 1\}^{(n+1)\ell}$ uniformly at random.
2. \mathcal{A} outputs m_0, m_1 .
3. Challenger calculates $c_{n-1} = \text{Enc}(k_{n-1}, m_1)$.
4. Challenger then gives (c_{n-1}, k_n, s_n) to \mathcal{A} .
5. \mathcal{A} returns \hat{b} .

Let $\Pr[H_2]$ be the chance of \mathcal{A} outputting 1.

Game 1: Finally circling back to the SSKeyLeak game, but now w/ bit $b = 1$.

1. Challenger samples $s_0 \leftarrow \{0, 1\}^\ell$.
2. Compute $(k_1, s_1, \dots, k_n, s_n) \leftarrow G^n(s_0)$.
3. \mathcal{A} outputs m_0, m_1 .
4. Challenger then calculates $c_{n-1} = \text{Enc}(k_{n-1}, m_1)$.
5. Give that value, (c_{n-1}, k_n, s_n) , to \mathcal{A} .
6. \mathcal{A} returns \hat{b} .

Let $\Pr[W_1]$ be the probability of \mathcal{A} outputting 1.

By the triangle inequality, the advantage is then bounded by the following expression:

$$|\Pr[W_0] - \Pr[W_1]| \leq |\Pr[W_0] - \Pr[H_1]| + |\Pr[H_1] - \Pr[H_2]| + |\Pr[H_2] - \Pr[W_1]|. \quad (1)$$

Reducing and Analyzing

Game 0 → Game H_1 Let \mathcal{B}_1 be a new adversary against G^n .

- \mathcal{B}_1 will receive a challenge string $r \in \{0,1\}^{(n+1)\ell}$ from G^n .
- \mathcal{B}_1 will then parse r as (k_1, \dots, k_n, s_n) .
- \mathcal{B}_1 will run \mathcal{A} to get m_0, m_1 .
- \mathcal{B}_1 then computes $c_{n-1} = \text{Enc}(k_{n-1}, m_0)$.
- Finally, \mathcal{B}_1 gives (c_{n-1}, k_n, s_n) to \mathcal{A} and its response.

Analyzing this:

- If r is outputted by $G^n(s_0)$, then this is just **Game 0**. So then $\Pr[\mathcal{B}_1 \rightarrow 1 \mid \text{Real}] = \Pr[W_0]$.
- In the same manner, if r is uniform random, then this is just **Hybrid 1**. Thus, $\Pr[\mathcal{B}_1 \rightarrow 1 \mid \text{Random}] = \Pr[H_1]$.

$|\Pr[W_0] - \Pr[H_1]| = \text{PRGadv}[\mathcal{B}_1, G^n]$, which is negligible.

Game $H_1 \rightarrow \text{Game } H_2$ \mathcal{B}_2 will now be the adversary against the semantic security of \mathcal{E} .

- \mathcal{B}_2 will run \mathcal{A} to get m_0, m_1 .
- \mathcal{B}_2 then sends m_0, m_1 to the SS challenger for \mathcal{E} .
- \mathcal{B}_2 will then get a ciphertext, c^* , from the challenger (which is just $\text{Enc}(k^*, m_b)$ for a rnd key k^*).
- \mathcal{B}_2 then samples $(k_n, s_n) \leftarrow \{0,1\}^{2\ell}$ uniformly at random.
- And then \mathcal{B}_2 will give (c^*, k_n, s_n) to \mathcal{A} , and output its response.

Looking at this:

- The key for the SS game, c^* , is rnd and independant of the leak. This is also the same distribution as k_{n-1} back in **Hybrid 1** and **Hybrid 2**.
- So then, if the SS challenger chooses m_0 , then this is just **Hybrid 1**.
- If instead the SS challenger chooses m_1 , then it is just **Hybrid 2**.

$|\Pr[H_1] - \Pr[H_2]| = \text{SSadv}[\mathcal{B}_2, \mathcal{E}]$, which is negligible.

Game $H_2 \rightarrow \text{Game 1}$ \mathcal{B}_3 will now be challenging the PRG security of G^n .

- \mathcal{B}_3 gets a challenge string $r \in \{0, 1\}^{(n+1)\ell}$ from G^n .
- \mathcal{B}_3 will parse r as (k_1, \dots, k_n, s_n) .
- \mathcal{B}_3 then runs \mathcal{A} to get m_0, m_1 .
- \mathcal{B}_3 computes $c_{n-1} = \text{Enc}(k_{n-1}, m_1)$.
- And finally, \mathcal{B}_3 gives (c_{n-1}, k_n, s_n) to \mathcal{A} and outputs its response.

Looking at this game:

- If r is outputted by $G^n(s_0)$, then this is just **Game 1**. This basically means that $\Pr[\mathcal{B}_3 \rightarrow 1 \mid \text{Real}] = \Pr[W_1]$.
- If instead r is uniform random, then this is just **Hybrid 2**, meaning that $\Pr[\mathcal{B}_3 \rightarrow 1 \mid \text{Random}] = \Pr[H_2]$.

$|\Pr[H_2] - \Pr[W_1]| = \text{PRGadv}[\mathcal{B}_3, G^n]$, which is negligible.

End

Combining the bounds, we get:

$$\text{SSKeyLeakadv}[\mathcal{A}, \mathcal{E}^1] \leq \text{PRGadv}[\mathcal{B}_1, G^n] + \text{SSadv}[\mathcal{B}_2, \mathcal{E}] + \text{PRGadv}[\mathcal{B}_3, G^n]. \quad (2)$$

I've already assumed, in the beginning, that G^n was secure, making the PRG advantages negligible. Also, since \mathcal{E} is semantically secure, then that also means that SS advantage is negligible. And since the sum of negligibles is negligible too, then $\therefore \text{SSKeyLeakadv}[\mathcal{A}, \mathcal{E}^1]$ is negligible.

Problem 2. (50 points at 10 each) Let $p > 2$ be an odd prime. For this problem, suppose the order of g is $p - 1$ (which is even). Such a g is called a *generator* of \mathbb{Z}_p^* . Such a g exists for every p (we won't prove this fact; see, for example, Section 7.5 of Victor Shoup's book at <http://www.shoup.net/ntb/ntb-v2.pdf> for the existence proof and Section 11.1 for how to sample it efficiently).

(a) Let $a = g^x$ for some $x \in \mathbb{Z}_{p-1}$ (the exponent works modulo $p - 1$ due to Fermat's little theorem). We can talk about the exponent x being even since $p - 1$ is even, and therefore $x + k(p - 1)$ has the same evenness as x for any $k \in \mathbb{Z}$.

Clearly if x is even, then a has a square root $g^{x/2}$ modulo p . Now show the converse: if a has a square root modulo p , then x is even. (*Hint: you can rewrite any element $b \in \mathbb{Z}_p^*$ as $b = g^y$ for some $y \in \mathbb{Z}_{p-1}$.*)

Solution. Your solution goes here

It'd be nice to check if a value is a square modulo p without having to explicitly know what power of g it is. Luckily we have the following test: a is a square iff $a^{(p-1)/2} \equiv_p 1$.

(b) Prove the forward direction: if a is square then $a^{(p-1)/2} \equiv_p 1$.

Solution. Your solution goes here

(c) Now prove the converse: if $a^{(p-1)/2} \equiv_p 1$, then a is square. (*Hint: Assume not and find a contradiction. Begin by writing a as g^x .*)

Solution. Your solution goes here

We have thus shown that exactly half the values in \mathbb{Z}_p^* have square roots, and we know how to identify them: by raising to $(p - 1)/2$. Note also that values that do have square roots have at least two of them: if $r \in \mathbb{Z}_p^*$ is a square root of $a \in \mathbb{Z}_p^*$, then so is $-r$, because $(-r)^2 = (-1) \cdot r \cdot (-1) \cdot r = (-1)(-1)r^2 = 1 \cdot r^2 = r^2$. The two square roots are different, because $p - r \neq r$ as p is odd. Therefore, each square cannot have more than two square roots by a simple counting argument: if some square had more than two square roots, there wouldn't be enough square roots for all the $(p - 1)/2$ squares, because just two square roots per square already takes up all the $p - 1$ possible square root values in \mathbb{Z}_p^* . Thus, each square has exactly two square roots.

(d) Show that if $(g^x)^2 \equiv_p a$, then $(g^{x+(p-1)/2})^2 \equiv_p a$, as well. Show that these two square roots are distinct: that is, show that $g^x \not\equiv_p g^{x+(p-1)/2}$.

Solution. Your solution goes here

We know from the paragraph above that a has only two square roots. But we have three values that all when squared give us a : g^x , $g^{x+(p-1)/2}$, and $-g^x$. Thus, it must be that $g^{x+(p-1)/2} \equiv_p -g^x$.

(e) Given that $g^{x+(p-1)/2} \equiv_p -g^x$, show that that $g^{(p-1)/2} \equiv_p -1$. Now show that for any $b \in \mathbb{Z}_p^*$ that is a non-square, $b^{(p-1)/2} \equiv_p -1$.

This refines our previous test for squares from parts (b) and (c): to test if something is a square, you raise it to $(p - 1)/2$ modulo p and check if the result is 1 or -1 . By the way, the value of $a^{(p-1)/2} \% p$ is called the Legendre symbol of a and is often written as $\left(\frac{a}{p}\right)$. The Legendre symbol can be generalized for composite p and this generalization is called the Jacobi symbol (we will not cover it here).

Solution. Your solution goes here